

Comp 206
Introduction to Operating Systems
Threads & Concurrency

Dr. Gikaru
Computer Science
Egerton University

Threads


- **Processes have the following components:**
 - ❖ an address space
 - ❖ a collection of operating system state
 - ❖ a CPU context ... or *thread* of control

- **On multiprocessor systems, with several CPUs, it would make sense for a process to have several CPU contexts (threads of control)**
 - ❖ Thread fork creates new thread not memory space
 - ❖ Multiple threads of control could run in the same memory space on a single CPU system too!

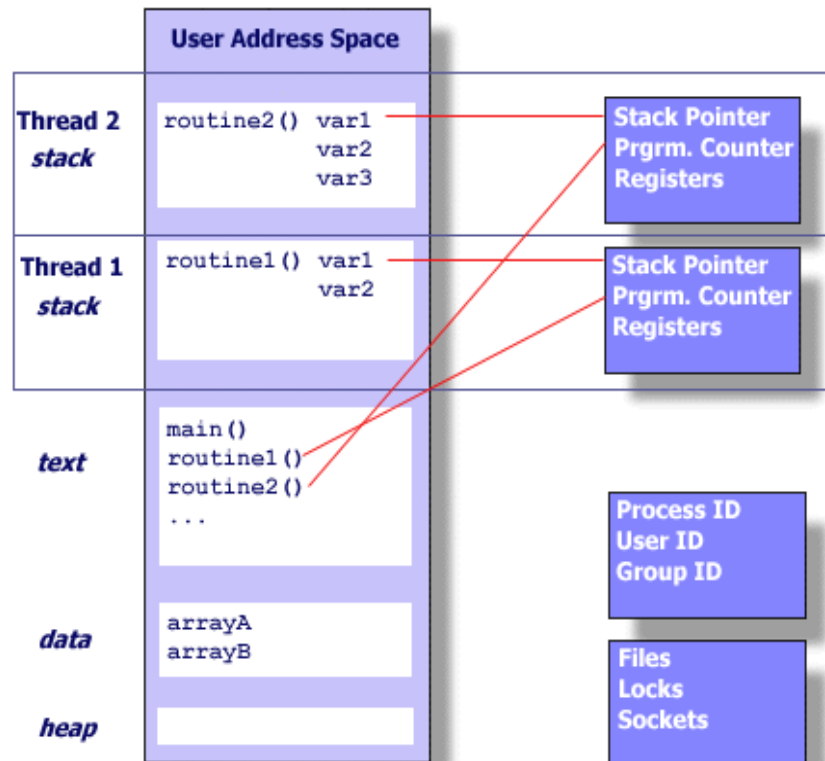
Threads

- ❑ Threads share a process address space with zero or more other threads
- ❑ Threads have their own
 - ❖ PC, SP, register state, stack
- ❑ A traditional process can be viewed as a memory address space with a single thread

Single thread state within a process



Multiple threads in an address space



What is a thread?

- **A thread executes a stream of instructions**
 - ❖ it is an abstraction for control-flow
- **Practically, it is a processor context and stack**
 - ❖ Allocated a CPU by a scheduler
 - ❖ Executes in the context of a memory address space

Summary of private per-thread state

Things that define the state of a particular flow of control in an executing program:

- ❖ Stack (local variables)
- ❖ Stack pointer
- ❖ Registers
- ❖ Scheduling properties (i.e., priority)

Shared state among threads

Things that relate to an instance of an executing program (that may have multiple threads)

- ❖ User ID, group ID, process ID
- ❖ Address space
 - Text
 - Data (off-stack global variables)
 - Heap (dynamic data)
- ❖ Open files, sockets, locks

Important: Changes made to shared state by one thread will be visible to the others

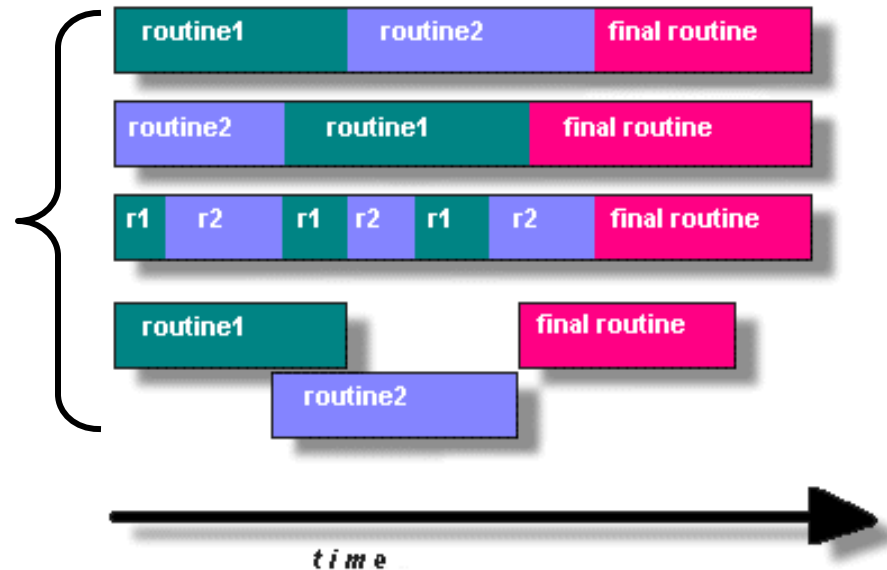
- ❖ Reading and writing memory locations requires synchronization! ... a major topic for later ...

How do you program using threads?

Split program into routines to execute in parallel

- ❖ True or pseudo (interleaved) parallelism


Alternative
strategies for
executing multiple
routines



Why program using threads?

- ❑ Utilize multiple CPU's concurrently
- ❑ Low cost communication via shared memory
- ❑ Overlap computation and blocking on a single CPU
 - ❖ Blocking due to I/O
 - ❖ Computation and communication
- ❑ Handle asynchronous events


Thread usage



A word processor with three threads


Processes versus threads - example

- **A WWW process**



Processes versus threads - example


- **A WWW process**



Why is this not a good web server design?


Processes versus threads - example

- **A WWW process**




Processes versus threads - example

- A WWW process

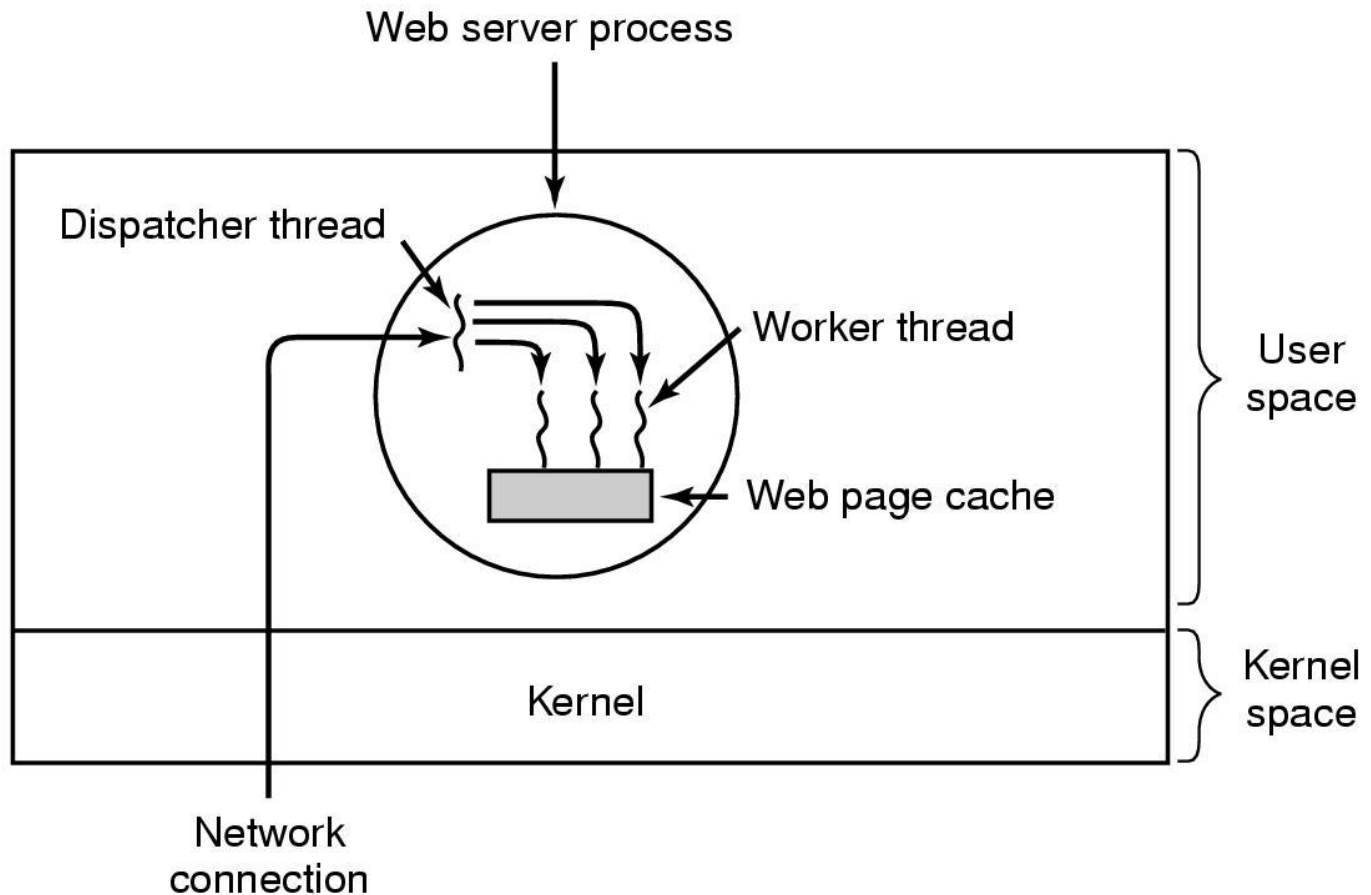


Processes versus threads - example

□ A WWW process



Threads in a web server



A multithreaded web server

Thread usage

```
while (TRUE) {  
    get_next_request(&buf);  
    handoff_work(&buf);  
}
```

(a)

```
while (TRUE) {  
    wait_for_work(&buf)  
    look_for_page_in_cache(&buf, &page);  
    if (page_not_in_cache(&page)  
        read_page_from_disk(&buf, &page);  
    return_page(&page);  
}
```

(b)

- **Rough outline of code for previous slide**
 - (a) Dispatcher thread
 - (b) Worker thread

System structuring options

Model	Characteristics
Threads	Parallelism, blocking system calls
Single-threaded process	No parallelism, blocking system calls
Finite-state machine	Parallelism, nonblocking system calls, interrupts

Three ways to construct a server

Common thread programming models

❑ **Manager/worker**

- ❖ Manager thread handles I/O and assigns work to worker threads
- ❖ Worker threads may be created dynamically, or allocated from a thread-pool

❑ **Pipeline**

- ❖ Each thread handles a different stage of an assembly line
- ❖ Threads hand work off to each other in a producer-consumer relationship


What does a typical thread API look like?

- ❑ POSIX standard threads (Pthreads)
- ❑ First thread exists in `main()`, typically creates the others
- ❑ **`pthread_create (thread, attr, start_routine, arg)`**
 - ❖ Returns new thread ID in "thread"
 - ❖ Executes routine specified by "start_routine" with argument specified by "arg"
 - ❖ Exits on return from routine or when told explicitly

Thread API (continued)

- ❑ **pthread_exit (status)**
 - ❖ Terminates the thread and returns "status" to any joining thread
- ❑ **pthread_join (threadid, status)**
 - ❖ Blocks the calling thread until thread specified by "threadid" terminates
 - ❖ Return status from pthread_exit is passed in "status"
 - ❖ One way of synchronizing between threads
- ❑ **pthread_yield ()**
 - ❖ Thread gives up the CPU and enters the run queue

Using create, join and exit primitives



An example Pthreads program

```
#include <pthread.h>
#include <stdio.h>
#define NUM_THREADS 5

void *PrintHello(void *threadid)
{
    printf("\n%d: Hello World!\n", threadid);
    pthread_exit(NULL);
}

int main (int argc, char *argv[])
{
    pthread_t threads[NUM_THREADS];
    int rc, t;
    for(t=0; t<NUM_THREADS; t++)
    {
        printf("Creating thread %d\n", t);
        rc = pthread_create(&threads[t], NULL, PrintHello, (void *)t);
        if (rc)
        {
            printf("ERROR: return code from pthread_create() is %d\n", rc);
            exit(-1);
        }
    }
    pthread_exit(NULL);
}
```

Program Output

```
Creating thread 0
Creating thread 1
0: Hello World!
1: Hello World!
Creating thread 2
Creating thread 3
2: Hello World!
3: Hello World!
Creating thread 4
4: Hello World!
```

For more examples see: <http://www.llnl.gov/computing/tutorials/pthreads>

Pros & cons of threads

□ Pros

- ❖ Overlap I/O with computation!
- ❖ Cheaper context switches
- ❖ Better mapping to shared memory multiprocessors

□ Cons


- ❖ Potential thread interactions
- ❖ Complexity of debugging
- ❖ Complexity of multi-threaded programming
- ❖ Backwards compatibility with existing code

User-level threads

- ❑ The idea of managing multiple abstract program counters above a single real one can be implemented using privileged or non-privileged code.
 - ❖ Threads can be implemented in the OS or at user level
- ❑ **User level thread implementations**
 - ❖ thread scheduler runs as user code (thread library)
 - ❖ manages thread contexts in user space
 - ❖ The underlying OS sees only a traditional process above


Kernel-level threads

The thread-switching code is in the kernel



User-level threads package

The thread-switching code is in user space



User-level threads

❑ Advantages

- ❖ cheap context switch costs among threads in the same process!
 - A procedure call not a system call!
- ❖ User-programmable scheduling policy

❑ Disadvantages

- ❖ How to deal with blocking system calls!
- ❖ How to overlap I/O and computation!

END